

Course progress

- ▶ Previously
 - ▶ PCA
 - ▶ OLS, ridge regression, conjugate gradient
 - ▶ Convex optimization, linear programming, Lagrange multipliers, duality and minimax games
 - ▶ Sparse regression (LASSO); NMF, Sparse PCA,
 - ▶ Dual ascent, dual decomposition, augmented Lagrangians, ADMM
 - ▶ Random sampling and randomized QR and SVD factorizations
 - ▶ Compressed sensing and matrix completion
 - ▶ DFT and FFT, shift invariant and circulant matrices/2D Fourier transform/filters
 - ▶ Graphs and their matrix representation, clustering
 - ▶ *Stochastic gradient descent, classification models and neural networks, CNNs*
- ▶ Today: Backpropagation and hyperparameters (Sec. VII.3 and VII.4)

NN models

- ▶ Training data $(v_1, y_1) \dots (v_n, y_n)$ for $v_i \in \mathbb{R}^m$ and labels y_i .
- ▶ Minimize

$$L(x) = \frac{1}{n} \sum_{i=1}^n \ell_i(x)$$

- ▶ The loss function is evaluated, e.g., cross-entropy loss for k -category classification, is evaluated at i -th data point

$$\ell_i(x) = \sum_{j=1}^k y_{j,i} \log F_j(x, v_i)$$

where $F(x, v)$ is given by the neural network,

$$v_0 = v,$$

$$v_{k+1} = R(A_{k+1}v_k + b_{k+1})$$

$$F(x, v) = v_K$$

and $x = (A_{1:K}, b_{1:K})$ represents all of its parameters.

NN models - weights $A_{1:K}$

- ▶ *Fully connected NN*
- ▶ Learning images
 - ▶ *CNN*
- ▶ Learning sequences
 - ▶ RNN, LSTM, GRU, Transformers
- ▶ Learning graphs
 - ▶ GNN

Gradient-based learning

- ▶ To learn x , we can follow a gradient-based descent algorithm

$$x_{k+1} = x_k - s_k \nabla L_{B_k}(x_k)$$

where the gradient over batch B_k is

$$\nabla_{B_k} L(x_k) = \frac{1}{|B_k|} \sum_{i \in B_k} \nabla \ell_i(x_k)$$

- ▶ In the GD we take the gradient of the full batch, i.e., B_k is always all of the n samples
- ▶ Minibatch GD chooses a batch of size of $B < n$ uniformly at random
- ▶ In SGD, $|B| = 1$, and the algorithm chooses a single $i(k)$ at step k uniformly at random
- ▶ Alternatively, in practice, the data is randomly ordered and the algorithm goes through it sequentially batch by batch

Gradient computation

- By the univariate chain rule

$$\frac{\partial}{\partial x_j} \ell = \frac{\partial}{\partial F} \ell(F) \frac{\partial F}{\partial x_j}$$

- Therefore

$$\nabla \ell(x) = \frac{\partial}{\partial F} \ell(F(x)) \nabla F(x)$$

where $\ell(x) := \ell_i(x)$ also depends on i -th true label y_i and $F(x) := F(x, v_i)$ also depends on the i -th feature vector v_i

- But since we're not differentiating with respect to y_i or v_i , we just treat them as parameters.

Chain rule

- In our earlier fully connected NN example

$$v_2 = F(A_k, b_k, v_0) = A_2 v_1 + b_2 = A_2(R(A_1 v_0 + b_1)) + b_2$$

where R is a set of 4 Relu activation functions

- The parameters or weights are

$$x = (A_1, b_1, A_2, b_2)$$

- Thus, we need a multivariate version of the chain rule. In 1D:

$$\frac{d}{dx} F_3(F_2(F_1(x))) = \left(\frac{dF_3}{dx}(F_2(F_1(x))) \right) \left(\frac{dF_2}{dx}(F_1(x)) \right) \left(\frac{dF_1}{dx}(x) \right)$$




Figure: Fig VII.2 from [1]

Chain rule

- ▶ The chain rules in the multivariate setting is
 $D(f \circ g) = Df \circ Dg$, so

$$\frac{\partial}{\partial A}[R(Av + b)] = \frac{\partial R}{\partial u} \circ \frac{\partial}{\partial A}(Av + b)$$

- ▶ Here $\frac{\partial R}{\partial u}$ is just the Jacobian matrix

$$\begin{bmatrix} \frac{\partial R_1}{\partial u_1} & \cdots & \frac{\partial R_1}{\partial u_n} \\ \vdots & \cdots & \vdots \\ \frac{\partial R_m}{\partial u_1} & \cdots & \frac{\partial R_m}{\partial u_n} \end{bmatrix}$$

- ▶ For $R(u) = \text{Relu}(u) = \max(0, u)$, it is just the diagonal matrix

$$\frac{\partial R_i}{\partial u_i} = \begin{cases} 0 & \text{if } (Av + b)_i < 0 \\ 1 & \text{if } (Av + b)_i > 0 \end{cases}$$

- ▶ What is the derivative of $u = Av$ w/r/t the matrix A ?

Derivative of matrix vector product

- ▶ What is the derivative of $u = Av$ w/r/t the matrix A ?
- ▶ This is an order 3 tensor: since

$$u_t = \sum_{\ell} A_{t\ell} v_{\ell}$$

we have

$$\frac{\partial u_t}{\partial A_{jk}} = \frac{\partial}{\partial A_{jk}} \left(\sum_{\ell} A_{t\ell} v_{\ell} \right) = v_k \delta_{tj}$$

- ▶ Therefore,

$$\frac{\partial R_i}{\partial A_{jk}} = \sum_t \frac{\partial R_i}{\partial u_t} \frac{\partial u_t}{\partial A_{jk}} = \sum_t \frac{\partial R_i}{\partial u_t} v_k \delta_{tj} = \frac{\partial R_i}{\partial u_j} v_k$$

- ▶ In this fashion can automatically compute all the derivatives going backwards (called autodiff or backprop)

Chain rule - order of multiplication

- ▶ The associativity of matrix multiplication gives two choices to compute ABC : either $(AB)C$ or $A(BC)$
- ▶ It is easy to see that for square matrices

$M_1 M_2 w$ needs $N^3 + N^2$ multiplications

while

$M_1(M_2 w)$ needs $N^2 + N^2$ multiplications

(Item N^3 is a simplification as are in fact subcubic algorithms)

- ▶ So forward propagation

$(M_1 M_2) M_3) \dots M_L w$ needs $(L - 1)N^3 + N^2$ multiplications

while backwards

$M_1(M_2(\dots M_L w))$ needs LN^2 multiplications

(Item N^3 is a simplification as are in fact subcubic algorithms)

Chain rule - order of multiplication

- ▶ The associativity of matrix multiplication gives two choices to compute ABC : either $(AB)C$ or $A(BC)$
- ▶ Let's say $A \in \mathbb{R}^{m \times n}$ and $B \in \mathbb{R}^{n \times q}$ The first way $(AB)C$

$$\begin{aligned}AB &= (m \times n)(n \times q) \text{ has } mnp \text{ multiplications} \\(AB)C &= (m \times p)(p \times q) \text{ has } mpq \text{ multiplications} \\&\text{Total } mp(n + q)\end{aligned}$$

- ▶ The second way $A(BC)$

$$\begin{aligned}BC &= (n \times p)(p \times q) \text{ has } npq \text{ multiplications} \\A(BC) &= (m \times n)(n \times q) \text{ has } mnq \text{ multiplications} \\&\text{Total } nq(m + p)\end{aligned}$$

Chain rule - order of multiplication

- ▶ If C is a column vector, $q = 1$, so $A(BC)$ has $n(m + p)$ steps vs $mnp + mp$ number of steps in $(AB)C$
- ▶ More generally we compare

$$\frac{mp(n + q)}{mnpq} = \frac{1}{q} + \frac{1}{n} \quad \text{vs} \quad \frac{nq(m + p)}{mnpq} = \frac{1}{m} + \frac{1}{p}$$

Chain rule

- ▶ Let's go back our feedforward fully connected NN example

$$w = v_2 = A_2 v_1 + b_2 = A_2(w(A_1 v_0 + b_1)) + b_2$$

- ▶ $\frac{\partial w}{\partial A_2}$ are given by

$$\frac{\partial w_i}{\partial A_{2jk}} = v_{1,k} \delta_{ij}$$

- ▶ and $\frac{\partial w}{\partial b_2}$ are

$$\frac{\partial w_i}{\partial b_{2j}} = \delta_{ij}$$




Figure: Fig VII.2 from [1]

Chain rule

- From the previous slide

$$w = v_2 = A_2 v_1 + b_2 = A_2(R(A_1 v_0 + b_1)) + b_2$$

- By the chain rule

$$\frac{\partial w}{\partial A_1} = \frac{\partial A_2(R(A_1 v_0 + b_1))}{\partial A_1} = A_2(R'(A_1 v_0 + b_1)) \frac{\partial (A_1 v_0 + b_1)}{\partial A_1}$$




Figure: Fig VII.2 from [1]

Adjoint methods

- ▶ The above notion that backpropagation optimizes the computation of derivatives is exploited in other large scale optimization problems.
- ▶ Let's say we want to solve

$$Ev = b(p)$$

- ▶ The solution vector

$$v(p) = E^{-1}b(p)$$

also depends on p and let's say the Jacobian $\frac{\partial v}{\partial p}$ is $N \times M$

Adjoint methods

- ▶ Taking the derivatives of both sides

$$Ev(p) = b(p)$$

- ▶ gives

$$E \frac{\partial v}{\partial p_j} = \frac{\partial b}{\partial p_j} \text{ for } j = 1, \dots, M$$

so

$$\frac{\partial v}{\partial p} = E^{-1} \frac{\partial b}{\partial p}$$

it seems that we have M linear systems of size N which will be expensive to solve iteratively as we minimize some loss function

$$F(v(p))$$

w/r/t p .

Adjoint methods

- ▶ If $F(v) = c^T v$ is linear, then

$$\frac{\partial F}{\partial p} = \frac{\partial F}{\partial v} \frac{\partial v}{\partial p} = c^T E^{-1} \frac{\partial b}{\partial p}$$

- ▶ Here multiply a row vector is multiplied by an $N \times N$ matrix E^{-1} and then by an $N \times M$ matrix $\frac{\partial b}{\partial p}$
- ▶ So you want to compute $(c^T E^{-1}) \frac{\partial b}{\partial p}$.
- ▶ The first step is equivalent to solving the adjoint equation

$$c = E^T \lambda \Rightarrow c^T = \lambda^T E \Rightarrow \lambda^T = c^T E^{-1}$$



$$\frac{\partial F}{\partial p} = \lambda^T \frac{\partial b}{\partial p}$$

which entails multiplying a row vector is multiplied by an $N \times M$

Gradient-based learning

- ▶ To learn x , i.e. find x such that $L(x)$ is a (possibly local) minimum, follow a gradient-based descent algorithm

$$x_{k+1} = x_k - s_k \nabla L_{B_k}(x_k)$$

where the gradient over batch B_k is

$$\nabla_{B_k} L(x_k) = \frac{1}{|B_k|} \sum_{i \in B_k} \nabla \ell_i(x_k)$$

- ▶ In the GD we take the gradient of the full batch, i.e., B_k is always all of the n samples
- ▶ Minibatch GD chooses a batch of size of $B < n$ uniformly at random
- ▶ In SGD, $|B| = 1$, and the algorithm chooses a single $i(k)$ at step k uniformly at random

Learning rate

- ▶ Several ways to determine s_k depending on the algorithm
- ▶ For GD, previous guarantee applied for a fixed step size $s \leq 1/M$ where M is the Lipschitz constant of the gradient

$$\|\nabla f(x) - \nabla f(y)\| \leq M\|x - y\|$$

uniformly over the domain (or a closed subset of the domain that includes the initialization and the minimum)

- ▶ Equivalently, if f is C^2 the eigenvalues λ_i of the Hessian of f are $|\lambda_i| \leq M$ for all i uniformly in x
- ▶ Often M is not known, but we can extend the convergence guarantees to *exact line search*

$$s_k = \arg \min_{s \geq 0} f(x_k - s \nabla f(x_k))$$

- ▶ and *backtracking line search*, which entails iteratively reducing s_k until

$$f(x_{k+1}) \leq f(x_k) - \frac{1}{2}s_k \|\nabla f(x_k)\|^2$$

Gradient decent convergence - fixed step size

- ▶ The previous argument didn't assume that f is convex, so the GD was converging to a local minimum.
- ▶ Finding a global minimum requires random x_k or grid search, which requires $t = O(1/\epsilon^d)$ iterations to achieve $\|\nabla f(x_k)\|_2^2 \leq \epsilon$ for functions with Lipschitz continuous gradients and $x \in \mathbb{R}^d$
- ▶ In practice gradient-based methods work well for non-convex functions used in NN even though there are not theoretical convergence guarantees
- ▶ If f is convex, then $\nabla f(x^*) = 0$ at a global minimizer x^*
- ▶ Therefore the above argument guarantees convergence to the global optimum at the above rate

Gradient decent convergence - fixed step size

- ▶ If f is *strongly* convex, i.e, the eigenvalues λ_i of the Hessian H are also $0 < m \leq \lambda_i$ uniformly in x , the convergence $O((1 - \frac{m}{M})^k)$ for $0 < c < 1$.
- ▶ This means that a bound of

$$f(x_k) - f(x^*) \leq \epsilon$$

can be achieved using only $O(\log(1/\epsilon))$ iterations. See previous lecture on convex optimization

- ▶ This rate is called “linear convergence” for historic reasons (the error lies below a line on a log-linear plot of the error vs iteration number)
- ▶ Typical loss functions in NN are not strongly convex, e.g., Relu and softmax are convex but not strongly convex
- ▶ Adding an ℓ^2 regularization term will make them such and can improve convergence

Nesterov accelerated descent

- ▶ If f is convex (but not necessarily strongly convex) is the $t = O(1/\epsilon)$ “sublinear convergence” optimal?
- ▶ So called Nesterov accelerated descent

$$x_{k+1} = y_k - s \nabla f(y_k)$$

$$y_{k+1} = x_{k+1} + \beta_k (x_{k+1} - x_k)$$

achieves error of $O(1/t^2)$ after t iterations.

- ▶ Can use $s = 1/M$ and $\beta_k = (k-1)/(k+2)$
- ▶ So only needs $t = O(1/\sqrt{\epsilon})$ to get within ϵ of the solution
- ▶ It not straightforward to understand why this method works better
- ▶ One observation is that this is not a descent method, i.e., the steps may overshoot the minimum and oscillate around it, rather than converging from one direction.
- ▶ Used in practice to optimize convex and nonconvex function in ML

Second order methods: Newton's method

- ▶ A second order approximation of a convex C^2 function is

$$g(y) = f(x) + \nabla f(y - x) + \frac{1}{2}(y - x)^T \nabla^2 f(x)(y - x)$$

- ▶ If the Hessian is positive definite,


$$\arg \min_y g(y) = x - (\nabla^2 f(x))^{-1} \nabla f(x)$$

- ▶ This idea leads to Newton's method

$$x_{k+1} = x_k - (\nabla^2 f(x_k))^{-1} \nabla f(x_k)$$

which has quadratic convergence under certain assumptions

- ▶ Rarely used in DL since Hessian computation is more computationally expensive, and the convergence of NN is only guaranteed when you're close to an optimal point



SGD step size

- ▶ In previous lecture, we showed sublinear convergence of SGD with a fixed step size

$$\min_{1 \leq k \leq t} \mathbb{E} \|\nabla f(x_k)\|^2 \leq \frac{C}{\sqrt{t}}$$

- ▶ So the norm is below ϵ if

$$\frac{C}{\sqrt{t}} \leq \epsilon$$

- ▶ This is guaranteed for

$$\frac{C^2}{\epsilon^2} \leq t$$

- ▶ So SGD requires $t = O(1/\epsilon^2)$ iterations to achieve $\mathbb{E} \|\nabla f(x_k)\|_2^2 \leq \epsilon$.
- ▶ vs GD that requires $t = O(1/\epsilon)$

SGD step size

- ▶ In practice reduce the step size as the optimization continues ($s_k \sim \frac{c}{k}$ or $\frac{c}{\sqrt{k}}$)
 - ▶ Convergence with probability 1 guarantees - Robbins Siegmund theorem
 - ▶ When the learning rates decrease too quickly, the expectation of the estimate takes too long to approach the optimum
 - ▶ When the learning rates decrease too slowly, the variance of the SGD estimate reduces to slowly
- ▶ When the Hessian of the objective is PD, can speed up the convergence of the expectation, but this does not reduce the variance
- ▶ Even though the convergence rate is slower, the per iteration time of SGD is faster
- ▶ Use SGD when the training time is the bottleneck. See Leon Bottou, *SGD Tricks*

Cross-validation

- ▶ K-fold cross validation
- ▶ Randomly spilt the data into K sets (epochs).
- ▶ Use one set as a training set and $K - 1$ as test sets.
- ▶ Repeat the procedure for different learning rates
- ▶ Find the best training set and the best learning rate that minimize the test error

Batch normalization

- ▶ We want the data in every minibatch to have uniform mean and variance.
- ▶ Let v_1, \dots, v_B be the size of the minibatch
- ▶ $V_i = (v_i - \mu) / \sqrt{\sigma^2 + \epsilon}$ for small $\epsilon > 0$, and sample mean μ and variance σ^2 .
- ▶ The input is

$$y_i = \gamma V_i + \beta$$

where γ and β are trainable parameters.

Dropout and regularization

- ▶ To avoid overfitting, for each data point v_i randomly dropout connections in NN according to Bernoulli distribution
- ▶ Leads to fewer parameters in each case but the full architecture is still available on average.
- ▶ L1 and L2 regularization with coefficients determined by cross-validation can be also used




Figure: Fig VII.6 from [1])

Next steps

- ▶ RNNs, LSTMs, GRUs, Transformers
- ▶ GNNs
- ▶ Kernel methods
- ▶ Bandit problems and reinforcement learning