

Course progress

- ▶ Previously
 - ▶ PCA: dimensionality reduction (simple unsupervised learning)
 - ▶ OLS, ridge regression, conjugate gradient
 - ▶ Convex optimization, linear programming, Lagrange multipliers, duality and minimax games
 - ▶ Sparse regression (LASSO); NMF, Sparse PCA,
 - ▶ Gradient descent, dual ascent, dual decomposition, augmented Lagrangians, ADMM
 - ▶ Random sampling and randomized QR and SVD factorizations
 - ▶ Compressed sensing and matrix completion
 - ▶ DFT and FFT, shift invariant and circulant matrices/2D Fourier transform/filters
 - ▶ Graphs and their matrix representation, clustering
 - ▶ *Stochastic gradient descent, classification models and neural networks, CNNs, Backprop, hyperparameters*
- ▶ Today: RNNs, Vanishing gradient, LSTM, GRU

NN models

- ▶ Fully connected NN
- ▶ Learning images
 - ▶ CNN
- ▶ Learning sequences
 - ▶ *RNN, LSTM, GRU*
- ▶ Learning graphs
 - ▶ GNN

NN models

- ▶ For the labelled data $(v_1, y_1) \dots (v_n, y_n)$ for $v_i \in \mathbb{R}^m$ and labels y_i

$$L(x) = \frac{1}{n} \sum_{i=1}^n \ell_i(x)$$

- ▶ the loss function can be e.g., cross-entropy loss for k -category classification

$$\ell_i(x) = \sum_{j=1}^k y_{j,i} \log F_j(x, v_i)$$

where $F(x, v)$ is given by the neural network and x represent all of its parameters.

Gradient computation

- By the univariate chain rule

$$\frac{\partial}{\partial x_j} \ell = \frac{\partial}{\partial F} \ell(F) \frac{\partial F}{\partial x_j}$$

- Therefore

$$\nabla \ell(x) = \frac{\partial}{\partial F} \ell(F(x)) DF(x)$$

where $\ell(x) := \ell_i(x)$ also depends on i -th true label y_i and $F(x) := F(x, v_i)$ also depends on the i -th feature vector v_i

- But since we're not differentiating with respect to y_i or v_i , we just treat them as parameters.

Chain rule

- ▶ The chain rules in the multivariate setting is
 $D(f \circ g) = Df \circ Dg$, so

$$D_A F = \frac{\partial}{\partial A} [R(Av + b)] = \frac{\partial R}{\partial u} \circ \frac{\partial}{\partial A} (Av + b)$$

for $u = Av + b$

- ▶ Here $\frac{\partial R}{\partial u}$ is just the Jacobian matrix

$$\begin{bmatrix} \frac{\partial R_1}{\partial u_1} & \cdots & \frac{\partial R_1}{\partial u_n} \\ \vdots & \cdots & \vdots \\ \frac{\partial R_m}{\partial u_1} & \cdots & \frac{\partial R_m}{\partial u_n} \end{bmatrix}$$

- ▶ The derivative w/r/t the matrix A is an order 3 tensor given by

$$\frac{\partial R_i}{\partial A_{jk}} = \sum_t \frac{\partial R_i}{\partial u_t} \frac{\partial u_t}{\partial A_{jk}} = \sum_t \frac{\partial R_i}{\partial u_t} v_k \delta_{tj} = \frac{\partial R_i}{\partial u_j} v_k$$

- ▶ In this fashion can automatically compute all the derivatives going backwards (called autodiff or backprop)

Chain rule

- For the fully connected NN

$$F = v_2 = A_2 v_1 + b_2 = A_2(R(A_1 v_0 + b_1)) + b_2$$

- By the chain rule

$$\frac{\partial F}{\partial (A_1)_{jk}} = A_2 \frac{\partial R}{\partial A_{jk}} = A_2 \frac{\partial R}{\partial u_j} v_k$$




Figure: Fig VII.2 from [1]

- Train to predict the next element in the sequence




Figure: Fig 8.4.2 from [4]

RNN

- ▶ For sequences,

$$v = (v_1, \dots, v^T)$$

like text or speech, want the predictions at t to depend on predictions at $t - 1$ and earlier

- ▶ RNN is given by

$$h_t = \sigma(Av_t + Bh_{t-1} + b)$$


$$= \sigma(Av_t + B\sigma(Av_{t-1} + Bh_{t-2} + b) + b)$$

$$F_t(v) = w_t = \text{softmax}(Ch_t)$$

where the matrices A , B and C are shared across time.

- ▶ Use the latent variable h_t to approximate the history

$$p(w_t | v_{t-1}, \dots, v_1) \approx p(w_t | h_{t-1})$$



Deep RNN

- Stack the hidden layers

$$h_t^l = R(A^l v_t^l + B^l h_{t-1})$$

where the matrices A^l , B^l and V^l are shared across time but not layers.

RNN

- ▶ Train to predict the next element in the sequence
- ▶ Since the prediction $F_{t,i}$ is a function of the parameters [we can use the cross-entropy loss for example]

$$\begin{aligned} L_T(A, B, C, b) &= \sum_{j \in \text{sequences}} \sum_{t=1}^T \ell(F_{t,j}(A, B, C, b), y_{t,j}) \\ &= \left[- \sum_{j \in \text{sequences}} \sum_{t=1}^T \sum_{i \in \text{categories}} y_{t,j,i} \log(F_{t,j,i}) \right] \end{aligned}$$




Figure: Fig 8.4.2 from [4]

- ▶ Let's fix the sequence and omit indexing by j .
- ▶ Since h_T depends on h_{T-1} , and h_{T-1} depends on h_{T-2} etc.

$$\begin{aligned}\frac{\partial}{\partial B} \ell(F_T) &= \frac{\partial \ell}{\partial F_T} \sum_{t=1}^T \frac{\partial F_T}{\partial h_t} \frac{\partial h_t}{\partial [B]^t} \\ &= \frac{\partial \ell}{\partial F_T} \frac{\partial F_T}{\partial h_T} \sum_{t=1}^T \left(\prod_{s=t+1}^T \frac{\partial h_s}{\partial h_{s-1}} \right) \frac{\partial h_t}{\partial [B]^t}\end{aligned}$$

where

$$\frac{\partial h_s}{\partial h_{s-1}} = \text{diag}(\sigma'(a_t))B$$

and

$$a_t = Av_t + Bh_{t-1} + b$$

RNN

- ▶ Based on the previous slide

$$\frac{\partial F_T}{\partial h_t} = \frac{\partial F_T}{\partial h_T} \sum_{t=1}^T \left(\prod_{s=t+1}^T \frac{\partial h_s}{\partial h_{s-1}} \right)$$

where

$$\frac{\partial h_s}{\partial h_{s-1}} = \text{diag}(\sigma'(a_s))B$$

- ▶ Often, the nonlinear activation functions have derivatives between 0 and 1, e.g.

$$0 < \tanh'(x) \leq 1$$

- ▶ Thus,

$$\left\| \frac{\partial h_s}{\partial h_{s-1}} \right\| \leq \max_i \sigma'(a_s)_i \|B\| \leq \|B\|$$

and therefore

$$\left\| \frac{\partial F_T}{\partial h_t} \right\| = \left\| \frac{\partial F_T}{\partial h_T} \right\| \left(\prod_{s=t+1}^T \left\| \frac{\partial h_s}{\partial h_{s-1}} \right\| \right) \leq \left\| \frac{\partial F_T}{\partial h_T} \right\| \|B\|^{T-t}$$

Vanishing gradient

- Therefore,

$$\left\| \frac{\partial}{\partial B} \ell(F_T) \right\| \leq \left\| \frac{\partial \ell}{\partial F_T} \right\| \left\| \frac{\partial F_T}{\partial h_T} \right\| \sum_{t=1}^T \|B\|^{T-t} \left\| \frac{\partial h_t}{\partial [B]^t} \right\|$$

- If $\|B\| < 1$ during the training process, the gradient decreases
- You don't know whether this is because the learning has finished, or you can't propagate because of the vanishing gradient

Exploding gradient

- ▶ Actually less of a problem
- ▶ Won't confuse with the learning being finished
- ▶ Can be solved by gradient clipping - simply rescaled the gradient when it's norm exceeds a certain threshold
- ▶ Based on the practical observation that in SGD learning, you can't trust the norm of the gradient anyway (the direction of gradient is more reliable)

- Add linear shortcut connections

$$h_t = \sigma(Av_t + Bh_{t-1} + b) + \sum_{\tau=1}^{t-1} g_{\tau}[h_{\tau}]^t$$

- Then

$$\frac{\partial h_t}{\partial h_{t'}} = \prod_{s=1}^{t-t'} \frac{\partial h_{t-s+1}}{\partial h_{t-s}} + \sum_{\tau=1}^{t-1} \frac{\partial [h_{\tau}]^t}{\partial h_{t'}} \frac{\partial g_{\tau}[h_{\tau}]^t}{\partial [h_{\tau}]^t}$$

- This linear connection prevents vanishing gradient, but requires learning g'_{τ} s, which can be computationally expensive for long sequences.

RNN+sparse update gates

- Instead of identity mappings, use

$$h_t = u_t \odot \sigma(Av_t + Bh_{t-1} + b) + (1 - u_t) \odot h_{t-1}$$

where the sparse update gates are

$$u_t = \sigma_u(A_u v_t + B_u h_{t-1} + b_u) \in [0, 1]^d$$

- Part of the learning is when to use sparse connections.

LSTM

- ▶ The above ideas led to LSTMs:

$$h_t = o_t \odot c_t$$

- ▶ the output gate determines how much of the previous states and the inputs (cell gate c^t) are exposed.

$$c_t = f_t \odot c_{t-1} + i_t \odot \tanh(A_c v_t + B_c c_{t-1} + b_c)$$

- ▶ the forget gate f_t is determines how much of the previous state is propagated linearly.

$$f_t = \sigma(A_f v_t + B_f c_{t-1} + V_f h_{t-1} + b_f)$$

- ▶ Parameters (weights) are “long-term memory”
- ▶ Previous outputs are “short-term memory”

LSTM

- ▶ Connections go back far in time
- ▶ Linear connections avoids vanishing gradients
- ▶ Minimize exploding gradients by sparsity
- ▶ h encodes the entire history and leads to embedding of sequences in \mathbb{R}^d

- ▶ Connections go back far in time
- ▶ Sparse connections essentially compress history in h_t
- ▶ Instead compute

$$h_t = \sum_{t'=1}^t w_{t'}(x_{t'}, x_t, t, t') \odot \hat{h}_{t'}(x_{t'}, t')$$

where $w_{t'}$ and $\hat{h}_{t'}$ depends on the position rather than history.

$$w_{t'}(x_{t'}, x_t, t, t') = \frac{\exp(Q(x_t, t)K(x_{t'}, t'))}{\sum_{t'=1}^t \exp Q(x_t, t)K(x_{t'}, t'))}$$

- ▶ Q is “query” and K is “key”

Next steps

- ▶ GNNs
- ▶ Kernel methods
- ▶ Bandits, RL
- ▶ Project presentations/Final quiz